

# BERT as Embedding

---

## Task: Finding Semantic Textual Similarity

```
# Ignore all your warnings
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

# Loading Libraries
import datetime
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import f1_score
import pylab
import scipy.stats as stats
from scipy.stats import boxcox

import re
import pickle

from tqdm import tqdm
import os
from wordcloud import WordCloud
from matplotlib_venn import venn2
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from textblob import TextBlob

import nltk
from nltk.probability import FreqDist
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
import string
eng_stopwords = stopwords.words('english')
import gc
from bs4 import BeautifulSoup
from nltk.stem import PorterStemmer
from nltk.tokenize import word_tokenize
import tensorflow

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import OrdinalEncoder
```

```

from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import normalize, StandardScaler
from scipy import sparse as sp
from scipy.sparse import hstack

from sklearn.metrics import pairwise_distances
from sklearn.preprocessing import minmax_scale
from sklearn.metrics.pairwise import cosine_similarity

from transformers import BertModel, BertTokenizer, BertForSequenceClassification
from torch.utils.data import Dataset, DataLoader
import torch

```

## Workflow

1. Problem Statement and Dataset Description
2. Dataset Loading
3. Machine learning Formulation
4. Exploratory Data Analyses
5. Preprocessing
6. Feature engineering and Modeling (Finding Similarity)
7. Results and Conclusion

.....

## 1. Problem Statement

Given two paragraphs, quantify the degree of similarity between the two text-based on Semantic similarity. Semantic Textual Similarity (STS) assesses the degree to which two sentences are semantically equivalent to each other. The STS task is motivated by the observation that accurately modelling the meaning similarity of sentences is a foundational language understanding problem relevant to numerous applications including machine translation (MT), summarization, generation, question-answering (QA), short answer grading, semantic search.

STS is the assessment of pairs of sentences according to their degree of semantic similarity. The task involves producing real-valued similarity scores for sentence pairs.

### About Dataset

- The data contains a pair of paragraphs. These text paragraphs are randomly sampled from a raw dataset.
- Each pair of the sentence may or may not be semantically similar. The candidate is to predict a value between 0-1 indicating a degree of similarity between the pair of text paras.
- 1 means highly similar
- 0 means highly dissimilar

**Note:** The given dataset does not contain any label.

.....

## 2. Dataset Loading

```
# Loading Dataset
```

```
data = pd.read_csv("Text_Similarity_Dataset.csv", delimiter=',')  
data.sample(10)
```

Unique_ID		text1	text2
2088	2088	china had role in yukos split-up china lent ru...	yahoo celebrates a decade online yahoo one of...
2582	2582	school tribute for tv host carson more than 1 ...	plaid mp s cottage arson claim a plaid cymru m...
18	18	mobile multimedia slow to catch on there is no...	text messages aid disaster recovery text messa...
3878	3878	wenger steps up row arsene wenger has stepped ...	bellamy under new fire newcastle boss graeme s...
1586	1586	sundance to honour foreign films international...	paraguay novel wins us book prize a novel set ...
2276	2276	wada will appeal against ruling the world anti...	uk athletics agrees new kit deal uk athletics ...
558	558	iranian mps threaten mobile deal turkey s bigg...	rank set to sell off film unit leisure group...
3685	3685	fears raised over ballet future fewer children...	consumer concern over rfid tags consumers are ...
1267	1267	tv calls after carroll error spurs boss martin...	parmalat sues 45 banks over crash parmalat has...
3214	3214	the year search became personal the odds are t...	kennedy begins pre-election tour liberal democ...

.....

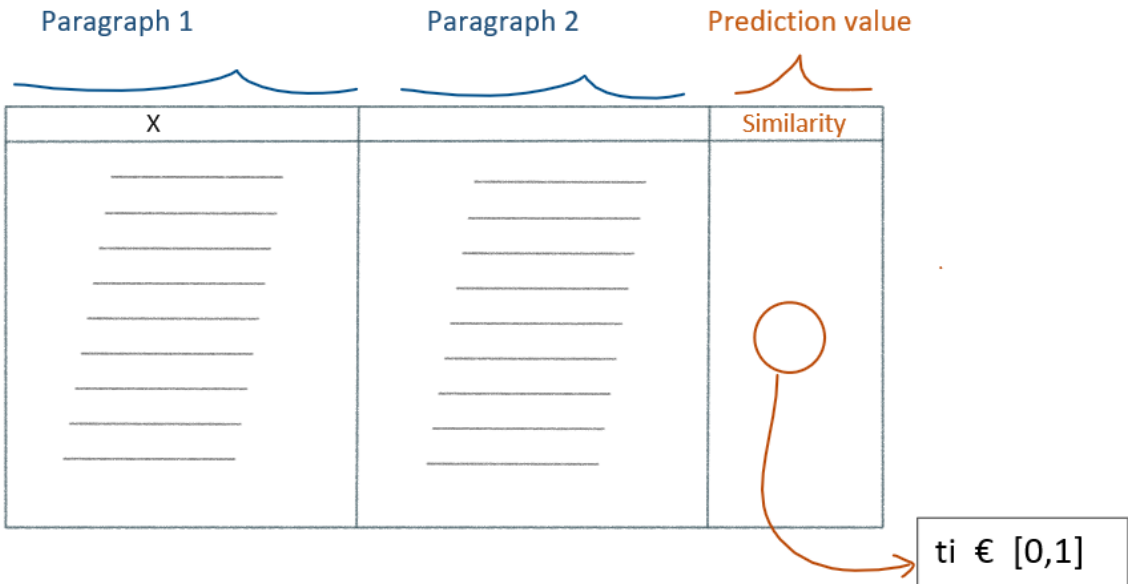
### 3. Machine learning Formulation

- The given dataset does not contain any label. Therefore, can be treated as an unsupervised learning problem.
- However, this does not imply that supervised techniques are not applicable.

So the task is to find the degree of similarity between the pair of text paras i.e predict a value between 0-1.

```
from IPython.display import Image  
Image(filename='Capture.png')
```

## Unsupervised Learning



### 3.1 Approaches to solve the Problem

### Approach 1: Using Avg Word2Vec of text1 and text2 and find similarity

1. Find the word embeddings of each word and average it on total words of text.
2. Find Cosine Similarity between embeddings of text1 and text2.
3. Scale the Cosine Similarity result between [0,1]

## Approach 2: Using Avg Tfidf Word2Vec of text1 and text2 and find similarity

1. Fit TfIDF on combination of text1 and text2 and create dictionary of vocabulary and idf\_scores
2. Find the word embeddings of each word and multiply the embedding of that word with its corresponding idf\_score and average the embeddings of total words present in text.
3. Following above step find the embeddings of text1 and text2.
4. Find Cosine Similarity between embeddings of text1 and text2.
5. Scale the Cosine Similarity result between [0,1]

### Approach 3: Using Spacy NLP pipeline to find similarity

1. Find the Similarity between embeddings of text1 and text2 using spacy NLP pipeline.
2. Scale the Similarity result between  $[0,1]$

### Approach 4: Using Bert pretrained model as feature extraction (pytorch framework GPU)

1. Format the text into desired format of Bert model.
2. Find the word 784 dim embeddings for each token and from the embedding of [cls] token.
3. Average the embeddings of total tokens passed to bert model.
4. Following above step find the embeddings of text1 and text2.
5. Find Cosine Similarity between embeddings of text1 and text2.
6. Scale the Cosine Similarity result between [0,1]

## 4. Exploratory Data Analyses

```
# Shape of Dataset
print(f"Shape of dataset: {data.shape}\n")

# Any Unique Values
print(f'Number of Duplicate values: {data.Unique_ID.duplicated().sum()}\n')

print(f" Any null Column{data.isnull().any()}")
```

Shape of dataset: (4023, 3)

Number of Duplicate values: 0

	Any null Column	Unique_ID	False
text1	False		
text2	False		

dtype: bool

### 4.0 Utility Function

```
# Utility function to plot lineplot and distplot using seaborn
def plot_sns(data, feature, color='lightblue', title=None, subtitle=None):

    """
    Utility function to plot lineplot and distplot using seaborn

    plot_sns(data, feature, color='lightblue', title=None, subtitle=None):

    data = data
    feature = coulum name
    color = color of plot
    title = Either 'length' or 'number' based on which to plot. Otherwise by de
    fault='None'
    subtitle = Either 'train_df' or 'test_df'. Otherwise by default='None'

    """
    f, (ax1, ax2) = plt.subplots(1, 2, figsize=(24, 6))

    # line plot
    sns.lineplot(np.arange(len(data)), data, ax=ax1, color=color)
    if title=='number':
        ax1.set(xlabel=f"Idx of {feature}", ylabel=f"Number of words in {featur
e}", title=f'Number of words in {feature} in {subtitle}\n')
    elif title=='length':
        ax1.set(xlabel=f"Idx of {feature}", ylabel=f"Length of {feature}", titl
e=f'Length of {feature} in {subtitle}\n')
    ax1.grid()

    # distribution plot
    sns.distplot(data, ax=ax2, color=color)
    if title=='number':
        ax2.set(xlabel=f"Idx of {feature}", ylabel=f"Number of words in {featur
e}", title=f'Number of words in {feature} in {subtitle}\n')
```

```

    elif title=='length':
        ax2.set(xlabel=f"Idx of {feature}", ylabel=f"Length of {feature}", title=f'Length of {feature} in {subtitle}\n')
        ax2.grid()
        plt.show()

#=====
#=====
# Utility function to plot bar graph for both train and test using seaborn
def plot_bar(train_data,test_data,feature=None,x_label=None, y_label=None):

    f, (ax1, ax2) = plt.subplots(1, 2, figsize=(24, 6))

    # for train_df
    sns.barplot(train_data,np.arange(len(train_data)),ax=ax1)
    ax1.set(xlabel=f"{x_label}", ylabel=f"{y_label} {feature}", title='train_df\n')
    ax1.grid()

    # for test_df
    sns.barplot(test_data,np.arange(len(test_data)),ax=ax2)
    ax2.set(xlabel=f"{x_label}", ylabel=f"{y_label} {feature}", title='test_df\n')
    ax2.grid()
    plt.show()

#=====
#=====
# Utility function to plot requency of most popular words
def word_frequency_plot(dataframe, title=None):
    list_of_all_words = []
    for sent in dataframe:
        list_of_all_words.extend(sent.split())

    top_50_words = pd.Series(list_of_all_words).value_counts()[:50]
    top_50_words_prob_dist = top_50_words.values/sum(top_50_words.values)

    # plot of frequency of polpular words in train
    plt.figure(figsize=(16,7))
    sns.barplot(top_50_words.index, top_50_words_prob_dist)
    plt.xlabel("words")
    plt.ylabel("frequency")
    plt.title(f"Frequency of most popular words {title}\n")
    plt.xticks(rotation=70)
    plt.grid()
    plt.show()

#=====
#=====
# Utility function to check if feature or variable follows Normal distribution
using Q-Q Plot
def q_q_plot(train_data, test_data, feature_name=None):

```

```

"""
# code refer: https://stackoverflow.com/a/13865874
"""

f, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 5))

measurements = train_data
stats.probplot(measurements, dist="norm", plot=ax1)
ax1.set(title=f'train : Q-Q Plot for {feature_name} \n')

measurements = test_data
stats.probplot(measurements, dist="norm", plot=ax2)
ax2.set(title=f'test : Q-Q Plot for {feature_name} \n')
plt.show()

#=====
#=====
#=====
# Utility function for box plot
def box_plot(data, feature_name=None):

    # for train data
    plt.figure(figsize=(26,4))
    sns.violinplot(data,color='darkred')
    plt.title(f'Train : violinplot Plot for {feature_name} \n')
    plt.xlabel(f"{feature_name}")
    plt.ylabel(f"Distribution")
    plt.grid()
    plt.show()

```

## 4.1. EDA: text1

### 4.1.1. Length of text1



```




# Length of text1
data['len_text1'] = data['text1'].apply(lambda x: len(x))

# plot
plot_sns(sorted(data['len_text1']), "text1", color='darkblue', title='length', subtitle='data')

# Box plot of Length of question_title in train and test
box_plot(sorted(data['len_text1']), "text1")

```





### Observation:


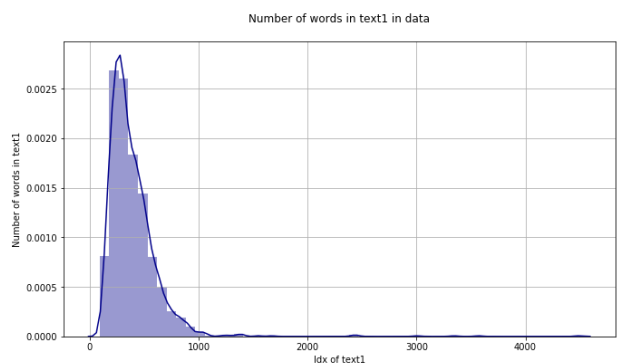
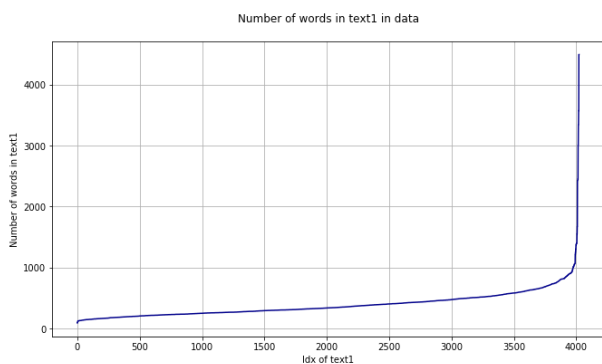
- Maximum length of text1 is upto 25000.
- Distribution is highly skewed toward right( looks like it following log normal distribution)

## 4.1.2. Number of words in text 1

```
# Number of words of text1
data['n_text1'] = data['text1'].apply(lambda x: len(x.split()))

# plot
plot_sns(sorted(data['n_text1']), "text1", color='darkblue', title='number', subtitle='data')

# Box plot of Length of question_title in train and test
box_plot(sorted(data['n_text1']), "number of words in text1")
```



### Observation:



- ### 4.1.3. WordCloud of text1

WordCloud of text1



## 4.2. EDA: text2




### 4.2.1. Length of text2

Length of text2 in data



Length of text2 in data








## 4.2.2. Number of words in text2

```
# Number of words of text1
data['n_text2'] = data['text2'].apply(lambda x: len(x.split(" ")))

# plot
plot_sns(sorted(data['n_text2']), "text2", color='darkblue', title='number', subtit
le='data')

# Box plot of Length of question_title in train and test
box_plot(sorted(data['n_text2']), "number of words in text2")
```



### Observation:

Observation is almost same as text1 observations.

## 4.2.3. WordCloud of text2


```
# refer: https://www.datacamp.com/community/tutorials/wordcloud-python

# For train_df
text_train = " ".join(word for word in data['text2'])

# Create and generate a word cloud image:
wordcloud = WordCloud().generate(text_train)

# Display the generated image:
plt.figure(figsize=(9,6))
plt.imshow(wordcloud, interpolation='bilinear')
plt.title("WordCloud of text2 \n")
plt.axis("off")
plt.show()
```

WordCloud of text2



## 5. Preprocessing

### Utility Function

```
# Preprocessing Functions
# credit : https://www.kaggle.com/urvishp80/quest-encoding-ensemble

mispell_dict = {"aren't" : "are not","can't" : "cannot","couldn't" : "could no
t","couldnt" : "could not","didn't" : "did not","doesn't" : "does not",
                "doesnt" : "does not","don't" : "do not","hadn't" : "had not",
"hasn't" : "has not","haven't" : "have not","havent" : "have not",
                "he'd" : "he would","he'll" : "he will","he's" : "he is","i'd"
: "I would","i'd" : "I had","i'll" : "I will","i'm" : "I am",
                "isn't" : "is not","it's" : "it is","it'll":"it will","i've" :
"I have","let's" : "let us","mightn't" : "might not",
                "mustn't" : "must not","shan't" : "shall not","she'd" : "she wo
uld","she'll" : "she will","she's" : "she is","shouldn't" : "should not",
                "shouldnt" : "should not","that's" : "that is","thats" : "that
is","there's" : "there is","theres" : "there is","they'd" : "they would",
                "they'll" : "they will","they're" : "they are","theyre": "they
are","they've" : "they have","we'd" : "we would","we're" : "we are",
                "weren't" : "were not","we've" : "we have","what'll" : "what wi
```

```

ll","what're": "what are","what's" : "what is","what've" : "what have",
        "where's" : "where is","who'd" : "who would","who'll" : "who wi
ll","who're" : "who are","who's" : "who is","who've" : "who have",
        "won't" : "will not","wouldn't" : "would not","you'd" : "you wo
uld","you'll" : "you will","you're" : "you are","you've" : "you have",
        "'re": " are","wasn't": "was not","we'll": " will","didn't": "di
d not","tryin'": "trying"}

```

```

imp_keywords = ("R", "r", "C", "c", "os", "OS")

```

```

# Counting the numeric feature and removing it

```

```

def count_digits_and_remove(text):

```

```

    """

```

```

    counting the number of occurance of digit

```

```

    return : text , (n_total_digit, n_2_digit, n_3_digit, n_4_digit, n_5_plus_d
igit)

```

```

    n_total_digit : Total occurance of numeric feature

```

```

    n_2_digit      : Number of time 2 digit numeric feature occur

```

```

    n_2_digit      : Number of time 3 digit numeric feature occur

```

```

    n_2_digit      : Number of time 4 digit numeric feature occur

```

```

    n_5_plus_digit : Number of time more than 4 digit numeric feature occur

```

```

    """

```

```

    digits = re.findall(r'[0-9]+',text)

```

```

    n_total_digit = []

```

```

    n_2_digit = []

```

```

    n_3_digit = []

```

```

    n_4_digit = []

```

```

    n_5_plus_digit = []

```

```

    n_total_digit.append(len(digits))

```

```

    for digit in digits:

```

```

        if len(digit)==2:

```

```

            n_2_digit.append(digit)

```

```

        elif len(digit)==3:

```

```

            n_3_digit.append(digit)

```

```

        elif len(digit)==4:

```

```

            n_4_digit.append(digit)

```

```

        elif len(digit)>4:

```

```

            n_5_plus_digit.append(digit)

```

```

    # remove all the numbers

```

```

    text = re.sub(r'[0-9]+'," ",text)

```

```

    return text , (len(n_total_digit), len(n_2_digit), len(n_3_digit), len(n_4_
digit), len(n_5_plus_digit))

```

```

#=====
# Counting number of non_alpha_numeric character and removing all the special character words
def count_non_alpha_numeric_and_remove(text):
    """ Counting number of non_alpha_numeric character(for programinng context) """

    # finding all the all the non_alpha_numeric char
    n_special_char = (re.findall(r"^[A-Za-z0-9 :]", text))

    # removing it.
    text = re.sub(r"^[A-Za-z0-9]", " ",text)

    return text ,len(n_special_char)

#=====
# Counting the number of all capital word(maybe it would be corelated with labels) and coverting into lower string character
def count_all_cap_words_and_lower_it(text):

    """Finding the number of all capital word and lower it"""

    # Find all the capital words
    n_all_capital_words = (re.findall(pattern = r'([A-Z]([A-Z]))+',string=text
))

    # converting into string into lower char string
    text = text.lower()

    return text , len(n_all_capital_words)

#=====
def _get_misspell(misspell_dict):
    misspell_re = re.compile('%s' % '|'.join(misspell_dict.keys()))
    return misspell_dict, misspell_re

def replace_typical_misspell(text):

    """De-Concatenation of words and correction of misspelled words"""
    misspellings, misspellings_re = _get_misspell(misspell_dict)

    def replace(match):
        return misspellings[match.group(0)]

    return misspellings_re.sub(replace, text)

#=====
# Return the number of links and text without html tags
# Also return the counts of 'number of lines' and remove it
def strip_html(text):

```

```

"""
Return the number of links and clean text (without html tags)
Also return the counts of 'number of lines' and remove it

"""
# finding http links using regex and counting it and remove it
n_links = (re.findall(r'http[s]?://\S+',text))
text = re.sub(r'http[s]?://\S+', " ",text)

# finding number of lines using regex and counting it and remove it
n_lines = re.findall(r'\n',text)
text = re.sub(r'\n', " ",text)

return text, len(n_links) , len(n_lines)

#=====
# function to remove all the stopwords and words having lengths less than 3
def remove_stop_words_and_punc(text) :

    """
    Remove all the stopwords
    Remove all the words whose length is less than 3 and not belong to important keywords (e.g. 'C','R','OS' etc)

    """
    # removing the words from the stop words list: 'no', 'nor', 'not'
    stops = set(stopwords.words("english"))
    stops.remove('no')
    stops.remove('nor')
    stops.remove('not')

    text= text.split()
    text = [w for w in text if not w in stops]

    # Removing the words having length less than 3 and not the imp_keyword
    clean_text = []
    for word in text:
        if word not in imp_keywords and len(word)<3:
            pass
        else:
            clean_text.append(word)

    clean_text = " ".join(clean_text)
    return(clean_text)

#=====
# function for stemming of words in text
def stem(text):
    stemmer = PorterStemmer()
    result = " ".join([ stemmer.stem(word) for word in text.split(" ")])
    return result

```

```

#=====
# Final text cleaning funtion
def clean_text(text, extra_features=True, strip_html_fn=True, count_all_cap_words_and_lower_it_fn=True, replace_typical_misspell_fn=True, count_digits_and_remove_fn=True, count_non_alpha_numeric_and_remove_fn=True, remove_stop_words_and_punc_fn=True, stem_fn=True):
    """
    This function sequentially execute all the cleaning and preprocessing function and finally gives cleaned text.
    Input: Boolean values of extra_features, strip_html, count_all_cap_words_and_lower_it, replace_typical_misspell, count_non_alpha_numeric_and_remove, remove_stop_words_and_punc, stem
           (by default all the input values = True)

    return: clean text

    """
    if strip_html_fn:
        # remove html tags
        clean_text, n_links, n_lines = strip_html(text)

    if count_all_cap_words_and_lower_it_fn:
        # Find all the capital words and covert all chars of text into lower char string
        clean_text, n_all_capital_words= count_all_cap_words_and_lower_it(clean_text)

    if replace_typical_misspell_fn:
        # de-concatenation of words
        clean_text = replace_typical_misspell(clean_text)

    if count_digits_and_remove_fn:
        # count the numbers and remove it
        clean_text, (n_total_digit, n_2_digit, n_3_digit, n_4_digit, n_5_plus_digit) = count_digits_and_remove(clean_text)

    if count_non_alpha_numeric_and_remove_fn:
        # Count the number of non alpha numeric character and remove it
        clean_text, n_non_alpha_char = count_non_alpha_numeric_and_remove(clean_text)

    if remove_stop_words_and_punc_fn:
        # removing Stopwords and the words length less than 3 (As these words mostly tend to redundant words) except 'C' and 'R' and 'OS' <-- programing keywords
        clean_text = remove_stop_words_and_punc(clean_text)

    if stem_fn:
        # stemming ( use only for BOW or TFIDF representation. Not effective for word embedding like w2v or glove)
        clean_text = stem(clean_text)

    # return extra_features

```

```

    if extra_features:
        return clean_text, (n_links, n_lines, n_all_capital_words, n_non_alpha_
char, n_total_digit, n_2_digit, n_3_digit, n_4_digit, n_5_plus_digit)

    else:
        return clean_text

#=====
#=====
#=====
# This function is only for word embedding pre processing
# This function will take dataframe and return cleaned dataframe (This finction
will only be used for BOW nd TFIDF representaion)
def clean_data_for_embeding(dataframe ):

    """
    This function will take dataframe and return cleaned dataframe along with e
xtra features.

    Input: dataframe which need to preprocess only for embeddings words
    Return: clean dataframe

    """
    cleaned_data = []

    for i in tqdm(range(dataframe.shape[0])):

        text = dataframe.iloc[i]

        cleaned_text = clean_text(text,extra_features= False, stem_fn=False) #
There is no need of extra_features to calculate again as it is already calcula
ted

        cleaned_data.append(cleaned_text)

    return cleaned_data

```

## 5.0. Preprocessing for word embedding (without stemming)

### Preprocessing utility function for word embedding (without stemming)

Only minor change in this function and above preprocessing function is this statement:

```
"clean_text(text,extra_features= False, stem_fn=False)" at 1
ine 19
```

## 5.1. Preprocessing: text1

```
clean_text_1_for_embedding = clean_data_for_embedding(data['text1'])
data['clean text1 for embedding'] = clean_text_1_for_embedding
```



## Sample

```
# Sample preprocessing
```

```
i=15
```

```
print(f"Before preprocessing: \n{'-'*20}\n{data['text1'][i]}\n")
```

```
print(f"After preprocessing: \n{'-'*20}\n{data['clean_text1_for_embedding'][i]}\n")
```

Before preprocessing:

-----

bond game fails to shake or stir for gaming fans the word goldeneye evokes excited memories not only of the james bond revival flick of 1995 but also the classic shoot-em-up that accompanied it and left n64 owners glued to their consoles for many an hour. adopting that hall of fame title somewhat backfires on this new game for it fails to deliver on the promise of its name and struggles to generate the original's massive sense of fun. this however is not a sequel nor does it relate to the goldeneye film. you are the eponymous renegade spy an agent who deserts to the bond world's extensive ranks of criminal masterminds after being deemed too brutal for mi6. your new commander-in-chief is the portly auric goldfinger last seen in 1964 but happily running around bent on world domination. with a determination to justify its name which is even less convincing than that of tina turner's similarly-titled theme song the game literally gives the player a golden eye following an injury which enables a degree of x-ray vision. rogue agent signals its intentions by featuring james bond initially and proceeding to kill him off within moments squashed by a plummeting helicopter. the notion is of course to add a novel dark edge to a 007 game but the premise simply does not get the juices flowing like it needs to. recent bond games like nightfire and everything or nothing were very competent and did a fine job of capturing the sense of flair invention and glamour of the film franchise. this title lacks that aura and when the bond magic shines through it feels like a lucky accident. the central problem is that the gameplay just is not good enough. quite aside from the bizarre inability to jump the even more bizarre glaring graphical bugs and dubious enemy ai the levels simply are not put together with much style or imagination. admittedly the competition has been tough even in recent weeks with the likes of halo 2 and half life 2 triumphing in virtually every department. what the game is good at is enveloping you in noisy dynamic scenes of violent chaos. as is the trend of late you are made to feel like you are in the midst of a really messy and fraught encounter. sadly that sense of action is outweighed by the difficulty of navigating and battling within the chaos meaning that frustration is often the outcome. and irregular save points mean you have to backtrack each time you are killed. a minute red dot passes for a crosshair although the collision-detection is so suspect that the difficulties of aiming weapons are compensated for. shooting enemies from a distance can be tricky and you will not always know you have picked them off since dead enemies vanish literally before they have fully hit the floor and they do so in some woefully uninspiring death animations. it is perhaps indicative of a lack of confidence that the game maker's allow you several different weapons almost immediately and throw you quickly into raging firefights - no time is risked with a measured build-up. by far the most satisfying element of the game is seeing old favourites like dr no goldfinger hat-friend oddjob and crazed russian sex beast xenia onatopp resurrected after all these years and with their faces rendered in an impressively recognisable fashion. there is a real thrill from doing battle with these legendary villains and it is a testament to the power of the bond universe that they can cut such a dash. but the in-game niggles combined with a story and presentation that just do not feel sufficiently well thought-through will make this a disappointment for most. diehard fans of bond will probably find enough here to make it a worthwhile purchase and try to ignore the failings. the game is weak not completely unplayable. then again 007 fanatics may also take umbrage at the cavalier blending of characters from different eras. given james bond's healthy pedigree in past games there is every reason to hope that this is just a blip a commendable idea that just has not worked that will be rectified when the character inevitably makes his return. goldeneye: rogue agent is out now

After preprocessing:

-----

bond game fails shake stir gaming fans word goldeneye evokes excited memories not james bond

revival flick also classic shoot accompanied left owners glued consoles many hour adopting h  
 allowed title somewhat backfires new game fails deliver promise name struggles generate orig  
 inal massive sense fun however not sequel nor relate goldeneye film eponymous renegade spy a  
 gent deserts bond world extensive ranks criminal masterminds deemed brutal new commander chi  
 ef portly auric goldfinger last seen happily running around bent world domination determinat  
 ion justify name even less convincing tina turner similarly titled theme song game literally  
 gives player golden eye following injury enables degree ray vision rogue agent signals inten  
 tions featuring james bond initially proceeding kill within moments squashed plummeting heli  
 copter notion course add novel dark edge game premise simply not get juices flowing like nee  
 ds recent bond games like nightfire everything nothing competent fine job capturing sense fl  
 air invention glamour film franchise title lacks aura bond magic shines feels like lucky acc  
 ident central problem gameplay not good enough quite aside bizarre inability jump even bizar  
 re glaring graphical bugs dubious enemy levels simply not put together much style imaginatio  
 n admittedly competition tough even recent weeks likes halo half life triumphing virtually e  
 very department game good enveloping noisy dynamic scenes violent chaos trend late made feel  
 like midst really messy fraught encounter sadly sense action outweighed difficulty navigatin  
 g battling within chaos meaning frustration often outcome irregular save points mean backtra  
 ck time killed minute red dot passes crosshair although collision detection suspect difficul  
 ties aiming weapons compensated shooting enemies distance tricky not always know picked sinc  
 e dead enemies vanish literally fully hit floor woefully uninspiring death animations perhap  
 s indicative lack confidence game maker allow several different weapons almost immediately t  
 hrow quickly raging firefights time risked measured build far satisfying element game seeing  
 old favourites like goldfinger hat fiend oddjob crazed russian sex beast xenia onatopp resur  
 rected years faces rendered impressively recognisable fashion real thrill battle legendary v  
 illains testament power bond universe cut dash game niggles combined story presentation not  
 feel sufficiently well thought make disappointment diehard fans bond probably find enough ma  
 ke worthwhile purchase try ignore failings game weak not completely unplayable fanatics may  
 also take umbrage cavalier blending characters different eras given james bond healthy pedig  
 ree past games every reason hope blip commendable idea not worked rectified character inevit  
 ably makes return goldeneye rogue agent



### 5.1.1. Number of words in clean\_text1\_for\_embedding


```

# Number of words of text1
data['n_word_clean_text1_for_embedding'] = data['clean_text1_for_embedding'].ap
ply(lambda x: len(x.split()))

# plot
plot_sns(sorted(data['n_word_clean_text1_for_embedding']), "clean_text1_for_embe
dding", color='darkblue', title='number', subtitle='data')

# Box plot of Length of question_title in train and test
box_plot(sorted(data['n_word_clean_text1_for_embedding']), "number of words in
clean_text1_for_embedding")
  
```





## 5.1.2. Distribution of number of words of text1 text before v/s after preprocessing

```
text_features_column = ['text1']

for idx,column in enumerate(text_features_column):

    # Calculating the length of text before and after preprocessing
    len_after_cleaning = data[f'clean_{column}_for_embedding'].apply(lambda x :
len(x.split()))
    len_before_cleaning = data[f'{column}'].apply(lambda x : len(x.split()))

    # plotting
    print( f"{idx+1}: Plot for {column}")
    plt.figure(figsize=(9,6))
    sns.distplot(len_before_cleaning, label=f'{column}')
    sns.distplot(len_after_cleaning, label=f'preprocessed_{column}')
    plt.title(f" Distribution of number of words of {column}  before v/s after
preprocessing\n",fontsize=15)
    plt.ylabel("distribtion")
    plt.xlabel(f"number of words in {column}")
    plt.legend()
    plt.grid()
    plt.show()
```

1: Plot for text1

Distribution of number of words of text1 before v/s after preprocessing





ers to send multimedia messages is really important for operators keen to squeeze more cash out of their customers and offset the cost of subsidising the handsets people are buying. the problem they face said shailendra jain head of mms firm adamind is educating people in how to send the multimedia messages using their funky handsets. also he said they have to simplify the interface so its not rocket science in terms of someone understanding it. research bears out the suspicion that people are not sending multimedia messages because they do not know how to. according to continental research 29% of the people it questioned said they were technophobes that tended to shy away from innovation. only 11% regarded themselves as technically savvy enough to send a picture or video message. the fact that multimedia services are not interoperable across networks and phones only adds to people's reluctance to start sending them said mr jain. they ask themselves: if i'm streaming video from one handset to another will it work he said. there's a lot of user apprehension about that. there are other deeper technical reasons why multimedia messages are not being pushed as strongly as they might. andrew bud executive chairman of messaging firm mblox said mobile phone operators cap the number of messages that can be circulating at any one time for fear of overwhelming the system. the rate we can send mms into the mobile network is fairly constant he said. the reason for this is that there are finite capacities for data traffic on the second generation networks that currently have the most users. no-one wants to take the risk of swamping these relatively narrow channels so the number of mms messages is capped said mr bud. this has led to operators finding other technologies particularly one known as wap-push to get multimedia to their customers. but when networks do find a good way to get multimedia to their customers the results can be dramatic. israeli technology firm celltick has found a way to broadcast data across phone networks in a way that does not overwhelm existing bandwidth. one of the first firms to use the celltick service is hutch india the largest mobile firm in the country. the broadcast system gets multimedia to customers via a rolling menu far faster than would be possible with other systems. while not multimedia messaging such a system gets people used to seeing their phones as a device that can handle all different types of content. as a result 40% of the subscribers to the hutch alive which uses celltick's broadcast technology regularly click for more pictures sounds and images from the operator. operators really need to start utilising this tool to reach their customers said yaron toren spokesman for celltick. until then multimedia will be a message that is not getting through.

After preprocessing:

-----

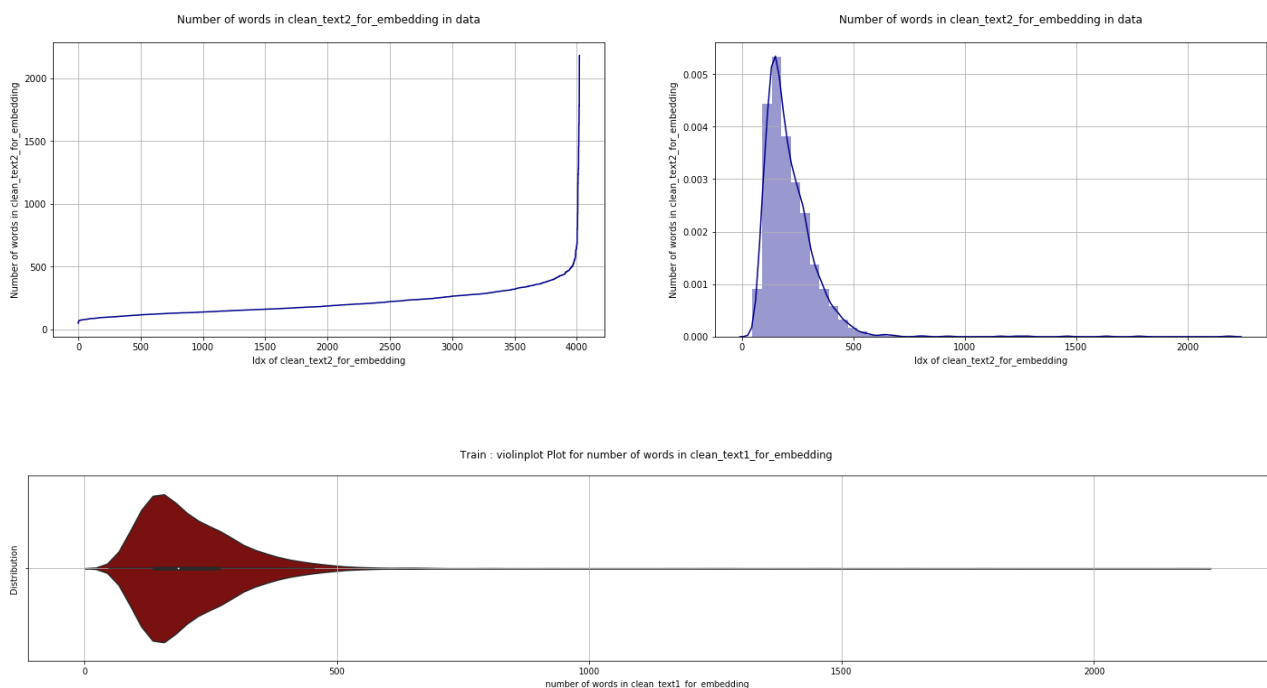
mobile multimedia slow catch doubt mobile phones sporting cameras colour screens hugely popular consumers swapping old phones slinkier dinkier versions thought responsible increase number phones sold third quarter according analysts gartner million handsets sold july september period according gartner analyst carolina milanesi seldom strong although consumers mobile's take send snaps sounds video clips far taking chance fact numbers people not taking sending pictures audio video growing figures gathered continental research shows british camera phone users never sent multimedia message mms despite fact period numbers camera phones doubled million getting mobile phone users send multimedia messages really important operators keen squeeze cash customers offset cost subsidising handsets people buying problem face said shailendra jain head mms firm adamind educating people send multimedia messages using funky handsets also said simplify interface not rocket science terms someone understanding research bears suspicion people not sending multimedia messages not know according continental research people questioned said technophobes tended shy away innovation regarded technically savvy enough send picture video message fact multimedia services not interoperable across networks phones adds people reluctance start sending said jain ask streaming video one handset another work said lot user apprehension deeper technical reasons multimedia messages not pushed strongly might andrew bud executive chairman messaging firm mblox said mobile phone operators cap number messages circulating one time fear overwhelming system rate send mms mobile network fairly constant said reason finite capacities data traffic second generation networks currently users one wants take risk swamping relatively narrow channels number mms messages capped said bud led operators finding technologies particularly one known wap push get multimedia customers networks find good way get multimedia customers results dramatic israeli technology firm celltick found way broadcast data across phone networks way not overwhelm existing bandwidth one first firms use celltick service hutch india largest mobile firm country broadcast system gets multimedia customers via rolling menu far faster would possible systems not multimedia messaging system gets people used seeing phones device handle different types content result subscribers hutch alive uses celltick broadcast technology regularly click picture

## 5.2.1. Number of words in clean\_text2\_for\_embedding

```
# Number of words of text1
data['n_word_clean_text2_for_embedding'] = data['clean_text2_for_embedding'].ap
ply(lambda x: len(x.split()))

# plot
plot_sns(sorted(data['n_word_clean_text2_for_embedding']), "clean_text2_for_embe
dding", color='darkblue', title='number', subtitle='data')

# Box plot of Length of question_title in train and test
box_plot(sorted(data['n_word_clean_text2_for_embedding']), "number of words in
clean_text1_for_embedding")
```



### Observation

- Preprocessing has reduced the large amount of number of words making distribution more symmetrical but it is still highly skewed towards right.
- We can see from above distribution that most of the text has less than 500 number of words. ( we can use 500 tokens as max\_len in BERT Embeddings)

## 5.2.2. Distribution of number of words of text1 text before v/s after preprocessing

```
text_features_column = ['text2']

for idx, column in enumerate(text_features_column):

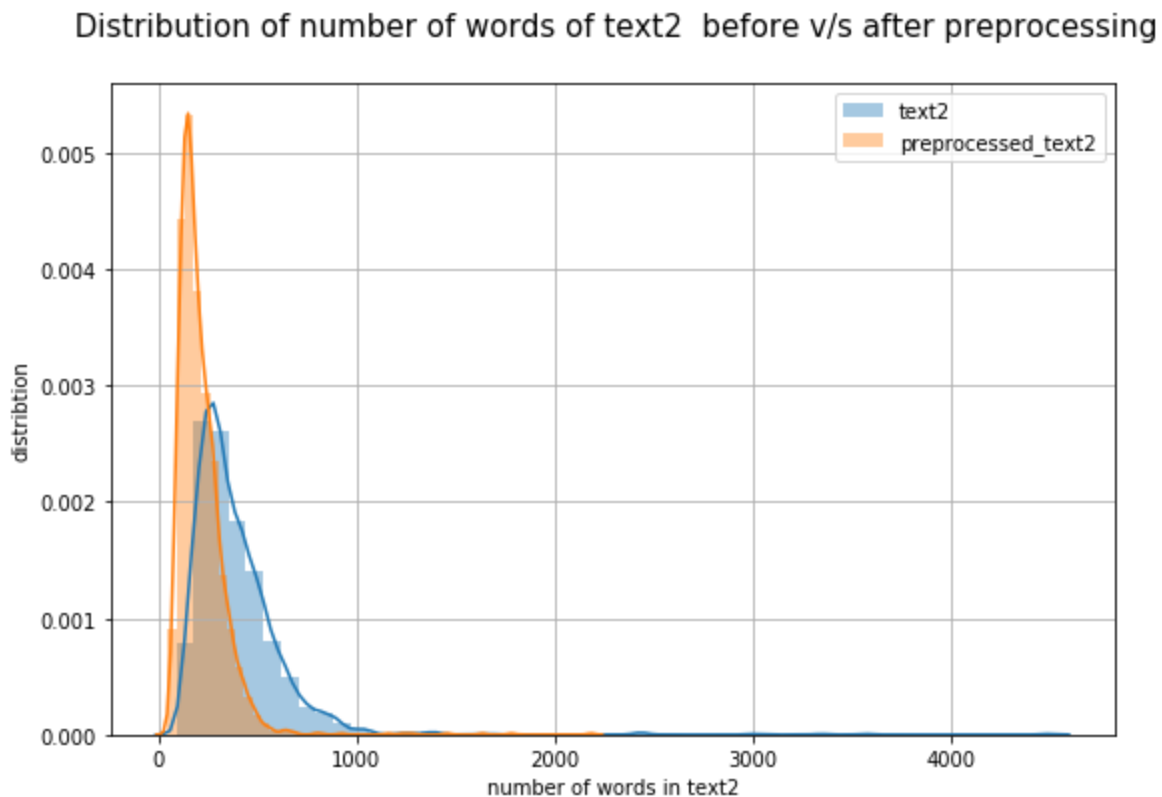
    # Calculating the length of text before and after preprocessing
    len_after_cleaning = data[f'clean_{column}_for_embedding'].apply(lambda x :
len(x.split()))
    len_before_cleaning = data[f'{column}'].apply(lambda x: len(x.split()))
```

```

# plotting
print( f"{idx+1}: Plot for {column}")
plt.figure(figsize=(9,6))
sns.distplot(len_before_cleaning, label=f'{column}')
sns.distplot(len_after_cleaning, label=f'preprocessed_{column}')
plt.title(f" Distribution of number of words of {column} before v/s after
preprocessing\n", fontsize=15)
plt.ylabel("distribtion")
plt.xlabel(f"number of words in {column}")
plt.legend()
plt.grid()
plt.show()

```

1: Plot for text2



## 6. Feature Engineering and Modeling (Finding Similarity)

### Vector Representation of text

Using count vectoriser and tfidf vectoriser directly make vocab size too large and might not good as well. Therefore I am skipping it.

#### 6.1. W2V Representation

#### 6.2. Tf-idf W2V Representation

#### 6.3. Using Bert Embedding

- Note: For word Embedding we have to preprocess the text again to remove stemming as it impact the sentiment of word in Embedding badly.

## 6.1. Approach 1: W2V Representation

Refer: <https://www.kaggle.com/phoenix9032/quest-preprocessing-data-for-embedding>

```
from gensim.models import KeyedVectors
```

```
news_path = 'crawl-300d-2M.vec'
```

```
embeddings_index = KeyedVectors.load_word2vec_format(news_path, binary=False)
```

```
## Building vocabulary from our Quest Data
```

```
def build_vocab(sentences, verbose = True):
```

```
    """
```

```
    :param sentences: list of list of words
```

```
    :return: dictionary of words and their count
```

```
    """
```

```
    vocab = {}
```

```
    for sentence in tqdm(sentences, disable = (not verbose)):
```

```
        for word in sentence:
```

```
            try:
```

```
                vocab[word] += 1
```

```
            except KeyError:
```

```
                vocab[word] = 1
```

```
    return vocab
```

```
=====
```

```
import operator
```

```
## This is a common function to check coverage between our quest data and the word embedding
```

```
def check_coverage(vocab, embeddings_index):
```

```
    a = {}
```

```
    oov = {}
```

```
    k = 0
```

```
    i = 0
```

```
    for word in tqdm(vocab):
```

```
        try:
```

```
            a[word] = embeddings_index[word]
```

```
            k += vocab[word]
```

```
        except:
```

```
            oov[word] = vocab[word]
```

```
            i += vocab[word]
```

```
        pass
```

```
    print('Found embeddings for {:.2%} of vocab'.format(len(a) / len(vocab)))
```

```
    print('Found embeddings for {:.2%} of all text'.format(k / (k + i)))
```

```
    sorted_x = sorted(oov.items(), key=operator.itemgetter(1))[:-1]
```

```
    return sorted_x
```

### 6.1.1. Check Coverage for clean\_text1\_for\_embedding

```
##Apply the vocab function to get the words and the corresponding counts
```



```
##Apply the vocab function to get the words and the corresponding counts
sentences = data["clean_text2_for_embedding"].apply(lambda x: x.split()).values
vocab = build_vocab(sentences)

print(f"\nFor clean_text2_for_embedding: \n{'-'*40}")
oov = check_coverage(vocab, embeddings_index)

## List 10 out of vocabulary word
print(f"\nTop 10 out of vocabulary word: \n{'-'*30}")
oov[:10]
```

```
[00:00<00:00, 483533.24it/s]
```

-----

### 6.1.3. Creating AVG W2V array

```
# For text1
cleaned_text1_avg_w2v = data['clean_text1_for_embedding'].apply(lambda x: avg_w2v_of_text(x))

# For text2
cleaned_text2_avg_w2v = data['clean_text2_for_embedding'].apply(lambda x: avg_w2v_of_text(x))
```

#### 6.1.4 Finding Similarity using W2V Representation

```
# Finding Similarity using pair wise distance
similarity_score_w2v=[]
for i in range(data.shape[0]):
    similarity_score_w2v.append(cosine_similarity([cleaned_text1_avg_w2v[i]], [cleaned_text2_avg_w2v[i]])[0][0])
```

```
# Scaling the similarity between [0-1]
similarity_score_w2v = minmax_scale(similarity_score_w2v, feature_range=(0, 1))

# Round off upto 2 decimal
similarity_score_w2v = np.round(similarity_score_w2v, decimals=3)
```

```
# Checking some similarity text
print(f"Index where similarity between text1 and text2 is greater than 0.9: {n
p.where(similarity_score_w2v>0.99) [0]}")

print(f"Similarity Score of those index: {similarity_score_w2v[np.where(similar
ity_score_w2v > 0.99)]}")
```

```
Index where similarity between text1 and text2 is greater than 0.9: [3403]
Similarity Score of those index: [1.]
```

```
# Sample of text where similarity score is greater than 0.99
i = 3403
print("Sample of text where similarity score is greater than 0.9\n\n")
print(f" text1: \n{'-'*7}\n{data.iloc[i].text1}\n")
print(f" text2: \n{'-'*7}\n{data.iloc[i].text2}\n")
```

```
Sample of text where similarity score is greater than 0.9
```

```
text1:
-----
holmes starts 2005 with gb events kelly holmes will start 2005 with a series of races in bri
tain. holmes will make her first track appearance on home soil since winning double olympic
gold in january s norwich union international in glasgow. she will also run in the grand pri
x in birmingham in february and may defend her indoor aaa 800m title in sheffield earlier th
at month. i am still competitive and still want to win she said. i m an athlete and i ca
n t wait to get back on the track. she added: these events are also a great opportunity to
thank the british public for the enormous levels of support they have given me from the mome
nt i stepped off that plane from greece. the glasgow meeting will see holmes compete over 1
500m in a five-way match against sweden france russia and italy.

text2:
-----
holmes starts 2005 with gb events kelly holmes will start 2005 with a series of races in bri
tain. holmes will make her first track appearance on home soil since winning double olympic
gold in january s norwich union international in glasgow. she will also run in the grand pri
x in birmingham in february and may defend her indoor aaa 800m title in sheffield earlier th
at month. i am still competitive and still want to win she said. i m an athlete and i ca
n t wait to get back on the track. she added: these events are also a great opportunity to
thank the british public for the enormous levels of support they have given me from the mome
nt i stepped off that plane from greece. the glasgow meeting will see holmes compete over 1
500m in a five-way match against sweden france russia and italy.
```

## 6.2. Approach 2: Tf-idf W2V Representation

```
def tfidf_w2v_of_dataframe(dataframe_text1, dataframe_text2):

    model = TfidfVectorizer()
    model.fit(dataframe_text1 + dataframe_text2)
```

```

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

# avg tfidf w2v conversion for text1
avg_tfidf_w2v_text1 = []
for i in tqdm(range(dataframe_text1.shape[0])):

    text = dataframe_text1.iloc[i]
    avg_tfidf_w2v_sent = np.zeros(300)
    len_of_text = len(text.split())
    weighted_sum = 0

    for word in text.split():

        try:
            # dictionary[word] = idf value of word in whole corpus
            # text.count(word)/len_of_text = tf value of word in this review

            idf_word = dictionary[word]
            tf_word = text.count(word)/len_of_text

            tf_idf_word = tf_word*idf_word
            weighted_sum += tf_idf_word

            avg_tfidf_w2v_sent += embeddings_index.word_vec(word) * tf_idf_

        except:
            pass

    avg_tfidf_w2v_sent = avg_tfidf_w2v_sent/weighted_sum
    avg_tfidf_w2v_text1.append(avg_tfidf_w2v_sent)

#=====

# avg tfidf w2v conversion for text2
avg_tfidf_w2v_text2 = []
for i in tqdm(range(dataframe_text2.shape[0])):

    text = dataframe_text2.iloc[i]
    avg_tfidf_w2v_sent = np.zeros(300)
    len_of_text = len(text.split())
    weighted_sum = 0

    for word in text.split():

        try:
            # dictionary[word] = idf value of word in whole corpus
            # text.count(word)/len_of_text = tf value of word in this review

            idf_word = dictionary[word]
            tf_word = text.count(word)/len_of_text

```



```
print(f" text1: \n{'-'*7}\n{data.iloc[i].text1}\n")
print(f" text2: \n{'-'*7}\n{data.iloc[i].text2}\n")
```

Sample of text where similarity score is greater than 0.9

text1:

-----

dvd copy protection strengthened dvds will be harder to copy thanks to new anti-piracy measures devised by copy protection firm macrovision. the pirated dvd market is enormous because current copy protection was hacked more than five years ago. macrovision says its new ripguard technology will thwart most but not all of the current dvd ripping (copying) programs used to pirate dvds. ripguard is designed to... reduce dvd ripping and the resulting supply of illegal peer to peer said the firm. macrovision said the new technology will work in nearly all current dvd players when applied to the discs but it did not specify how many machines could have a problem with ripguard. some bbc news website users have expressed concerns that the new technology will mean that dvds will not work on pcs running the operating system linux. the new technology will be welcomed by hollywood film studios which are increasingly relying on revenue from dvd sales. the film industry has stepped up efforts to fight dvd piracy in the last 12 months taking legal action against websites which offer pirated copies of dvd movies for download. ultimately we see ripguard dvd... evolving beyond anti-piracy and towards enablement of legitimate online transactions interoperability in tomorrow's digital home and the upcoming high-definition formats said steve weinstein executive vice president and general manager of macrovision's entertainment technologies group. macrovision said ripguard would also prevent against rent rip and return - where people would rent a dvd copy it and then return the original. ripguard is expected to be rolled out on dvds from the middle of 2005 the company said. the new system works specifically to block most ripping programs - if used those programs will now most likely crash the company said. macrovision has said that rip guard can be updated if hackers find a way around the new anti-copying measures.

text2:

-----

dvd copy protection strengthened dvds will be harder to copy thanks to new anti-piracy measures devised by copy protection firm macrovision. the pirated dvd market is enormous because current copy protection was hacked more than five years ago. macrovision says its new ripguard technology will thwart most but not all of the current dvd ripping (copying) programs used to pirate dvds. ripguard is designed to... reduce dvd ripping and the resulting supply of illegal peer to peer said the firm. macrovision said the new technology will work in nearly all current dvd players when applied to the discs but it did not specify how many machines could have a problem with ripguard. the new technology will be welcomed by hollywood film studios which are increasingly relying on revenue from dvd sales. the film industry has stepped up efforts to fight dvd piracy in the last 12 months taking legal action against websites which offer pirated copies of dvd movies for download. ultimately we see ripguard dvd... evolving beyond anti-piracy and towards enablement of legitimate online transactions interoperability in tomorrow's digital home and the upcoming high-definition formats said steve weinstein executive vice president and general manager of macrovision's entertainment technologies group. macrovision said ripguard was designed to plug the digital hole that was created by so-called decss ripper software. it circumvents content scrambling system measures placed on dvds and let people make perfect digital copies of copyrighted dvds in minutes. those copies could then be burned onto a blank dvd or uploaded for exchange to a peer-to-peer network. macrovision said ripguard would also prevent against rent rip and return - where people would rent a dvd copy it and then return the original. ripguard is expected to be rolled out on dvds from the middle of 2005 the company said. the new system works specifically to block most ripping programs - if used those programs will now most likely crash the company said. macrovision has said that rip guard can be updated if hackers find a way around the new anti-copying measures.

## 6.3. Approach 3: Finding Similarity Using Spacy api

```

import spacy
import en_core_web_sm

nlp = spacy.load('en_core_web_sm')

similarity_score_spacy = []
for i in range(data.shape[0]):

    doc1 = nlp(data['clean_text1_for_embedding'].iloc[i])
    doc2 = nlp(data['clean_text2_for_embedding'].iloc[i])
    similarity_score_spacy.append(doc1.similarity(doc2))

# Scaling the similarity between [0-1]
similarity_score_spacy = minmax_scale(similarity_score_spacy, feature_range=(0,
1))

# Round off upto 2 decimal
similarity_score_spacy = np.round(similarity_score_spacy, decimals=3)

```

```

# Checking some similarity text
print(f"Index where similarity bestween text1 and text2 is greater than 0.9: {n
p.where(similarity_score_spacy>0.95) [0]}")

print(f"Similarity Score of those index: {similarity_score_spacy[np.where(simil
arity_score_spacy>0.95)]}")

```

```

Index where similarity bestween text1 and text2 is greater than 0.9: [ 462  684 1483 2270 22
84 3403 3859]
Similarity Score of those index: [0.953 0.951 0.951 0.953 0.988 1.      0.951]

```

```

# Sample of text where similarity score is greater than 0.99
i = 2488

print("Sample of text where similarity score is greater than 0.9\n\n")
print(f" text1: \n{'-'*7}\n{data.iloc[i].text1}\n")
print(f" text2: \n{'-'*7}\n{data.iloc[i].text2}\n")

```

```

Sample of text where similarity score is greater than 0.9

```

```

text1:
-----

```

```

gadgets galore on show at fair the 2005 consumer electronics show in las vegas is a geek s p
aradise with more than 50 000 new gadgets and technologies launched during the four-day even
t. top gadgets at the show are highlighted in the innovations showcase which recognises so
me of the hottest developments in consumer electronics. the bbc news website took an early p
re-show look at some of those technologies that will be making their debut in 2005. one
of the key issues for keen gadget users is how to store all their digital images audio and
video files. the 2.5gb and 5gb circular pocket hard drive from seagate might help. the exter
nal usb drive won a ces best innovations design and engineering award and is small enough to
slip into a pocket. it is the kind of storage that appeals to people who want their pcs to
look cool said seagate. it is all about style but it also has lots of functionality. it
is the first time you can say a hard drive is sexy it said. in the centre of the device is
a blue light that flashes while data is being written to ensure users do not unplug it when

```

it is busy saving those precious pictures. universal electronics' nevos1 is a universal controller that lets people use one device to get at their multimedia content such as photos no matter where it is in their house. it can also act as a remote for home theatre and stereo systems. working with home broadband networks and pcs the gadget has built-in wireless and a colourful simple interface. paul arling uei chief said consumers face real problems when trying to get at all the files they own that are typically spread across several different devices. he said the nevo gave people a simple single way to regain some control over digital media in the home. the nevo won two awards at ces one as a girl's best friend award and another for innovation design and engineering. the gadget is expected to go on sale before the summer and will cost about \$799 (£425). hotseat is targeting keen gamers with money to spend with its solo chassis gaming chair. the specially-designed chair lets gamers play in surround-sound while stretching out in their own space. it is compatible with all the major games consoles dvd players and pcs. we found that kids love playing in surround sound said jay leboff from hotseat. we are looking at offering different types of seats depending on the market success of this one. the chair also lets people experience surround sound while watching videos with wireless control for six surround sound speakers. and a drinks holder. the chair which looks like a car seat on a skeletal frame should go on sale in april and is expected to cost \$399 (£211). satellite radio is big business in the us. in the uk the digital radio technology is known as dab and works on slightly different technology. eton corporation's porsche designed p7131 digital radio set will be launched both as a dab radio in the uk as well as a satellite radio set in the us. dab sets have been slow to take-off in the uk but this one concentrates on sleek looks as much as technology. it is for the risk-conscious consumer said an eton spokesperson. we are proud of it because it has the sound quality for the audiophile and the looks for the design-conscious consumer. the porsche radio is set to go on sale at the end of january in the us and in the first quarter of 2005 in the uk. in the us it is expected to cost \$250 (£133). the average person has a library of 600 digital images estimates the consumer electronics association the organisation behind ces. this is expected to grow to a massive 3 420 images - or 7.2gb - in five years time. one gadget that might help swell that collection is sanyo's tiny handheld vpc-c4 camcorder which is another innovation in design and engineering award winner. it combines high quality video and stills in a very small device. it takes mpeg4 video quality at 30 frames a second and has a four megapixel still camera. images and video are stored on sd cards which have come down in price in recent months. a 512mb card will store about 30 minutes of video and 420 stills. the device is so tiny it can be controlled with one thumb. because images and video are stored on sd memory it is portable to other devices and means other data like audio can be stored on the card too. wearable technology has always promised much but failed to deliver because of lack of storage capability and poor design. mpio's tiny digital usb music players come in an array of fashionable colours taking a leaf out of the apple ipod mini book of design and reflecting the desire for gadgets that look good. slung on a cord the player would not look too geeky dangling discreetly from the neck. although the pendant design was launched three months ago the device emphasises large storage as well as good looks for fashion-conscious gadget fiends. an even dinkier model the fy500 comes out in may and will store about 256mb of music. the range of players recently won an international forum design award 2005.

text2:

-----

rivals of the £400 apple... the mac mini is the cheapest apple computer ever. but though it is cheap for a mac how does it compare to pcs that cost about the same amount dot.life tries to find out if you can get more for your money if you stick with the beige box. an extremely small computer that is designed to bring the macintosh to the masses. apple offer a less powerful mac mini for £339 but the £399 model has a 1.4ghz power pc chip 80 gigabyte hard drive combined cd burner/dvd player. it comes equipped with usb and firewire ports for peripheral connections ethernet port for broadband a port for standard video output and an audio/headphone jack. the machine comes with mac os x the apple operating system the software suite ilife which includes itunes iphoto imovie idvd and garageband. a monitor keyboard or mouse. there is also no built-in support for wireless technology or any speakers. the lack of a dvd burner is an omission in the age of backing-up important software. wireless and a dvd burner can be added at extra cost. apple are targeting people who already have a main computer and want to upgrade - especially pc users who have used an apple ipod. compact and stylish the mac mini would not look out of place in any home. apple computers are famously user friendly and offer much better network security which means fewer viruses.



the package of software that comes with the machine is the best money can buy. the mac mini is just a box. if you don't already have a monitor etc adding them to the package sees the value for money begin to dwindle. macs don't offer the upgrade flexibility of a pc and the machine's specifications lack the horse power for tasks such as high-end video editing or games. the mac mini puts the macintosh within the reach of everyone an apple spokesman said. it will bring more customers to the platform especially pc users and owners. an entry-level machine designed for basic home use. a 2.6ghz intel celeron chip 40 gigabyte hard drive 256mb combined cd burner/dvd player. it comes equipped with a 17 inch monitor keyboard and mouse. the machine has 6 usb ports and an ethernet port for broadband connection. there's also a port for standard video output. the machine comes with windows xp home edition. it provides basic home tools such as a media player and word processor. a dvd burner or any wireless components built in. wireless and a dvd burner can be added at extra cost. homes and small offices including those looking to add a low cost second computer. cost is the clear advantage. the dell provides enough power and software for basic gaming and internet surfing. it's easily upgradeable so a bigger hard drive better sound and graphics cards can be added. the dell is hardly stylish and the hard drive is on the small size for anyone wanting to store photos or a decent sized digital music collection. this machine is for small businesses and for people who want a second computer for basic home use perhaps in a kids bedroom a spokesman for dell said. i think we offer better value once you realise all the extras needed for the mac mini. a desktop computer that pc pro magazine dubbed best performer in a group test of machines that cost only £399 (£469 including vat). a good basic pc that according to pc pro has superb upgrade potential. for your money you get a 1.8ghz amd sempron processor 512mb of ram 120gb hard drive dvd writer 16-inch monitor mouse keyboard and windows xp2 much more than the basics. it cannot handle 3d graphics and has no firewire slots. those on a limited budget who want a machine they can add to and improve as their cash allows. it's cheap and has plenty of room to improve but that could end up making it expensive in the long run. it's a good basic workhorse. it's not pretty and has a monitor rather than a flat-panel display. some of the upgrades offered by jal to the basic model are pricey. you might find that you want to chop and change quite quickly. nick ross deputy labs editor at pc pro said the important point about buying a cheap and cheerful pc is the upgrade path. interest has switched from processor power to graphics and sound cards as that's what makes the difference in games. even manufacturers are not going to be marketing machines as faster he said they'll emphasise the different features. a computer built from bits you buy and put together yourself. a surprisingly good pc sporting an amd athlon xp 2500 processor 512 megabytes ram a graphics card with 128 ram on board plus tv out a 40 gb hard drive cd-writer and dvd player windows xp home. anything else. you're building it so you have to buy all the software you want to install and do your own trouble-shooting and tech support. building your own machine is easier than it used to be but you need to read specifications carefully to make sure all parts work together. experienced and keen pc users. building your own pc or upgrading the one you have is a great way to improve your understanding of how it all works. it's cheap you can specify exactly what you want and you get the thrill of putting it together yourself. and a bigger thrill if everything works as it should. once it's built you won't be able to do much with it until you start buying software for it. if it starts to go wrong it might take a lot of fixing. as gavin cox of the excellent buildyourown.org.uk website put it: it will be tough to obtain/build a pc to ever be as compact and charming as the mac mini. performance-wise it's not cutting edge and is barely entry-level by today's market but up against the mac mini i believe it will hold its own and even pull a few more tricks says gavin cox. the good news is that the machine is eminently expandable. by contrast says mr cox the mac mini is almost disposable.

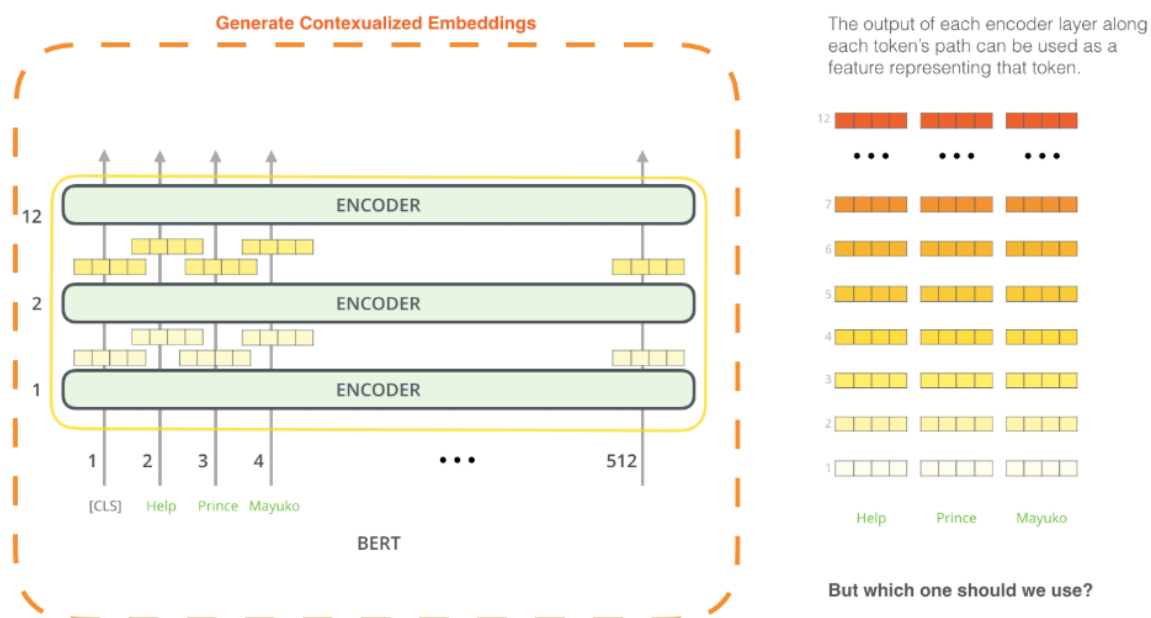
## 6.4. Approach 4: Using AVG BERT Embedding to find Similarity

Refer: <https://jalammar.github.io/illustrated-bert/>

```
from IPython.display import Image
Image(filename='bert_fine_tuning.png')
```

## BERT for feature extraction

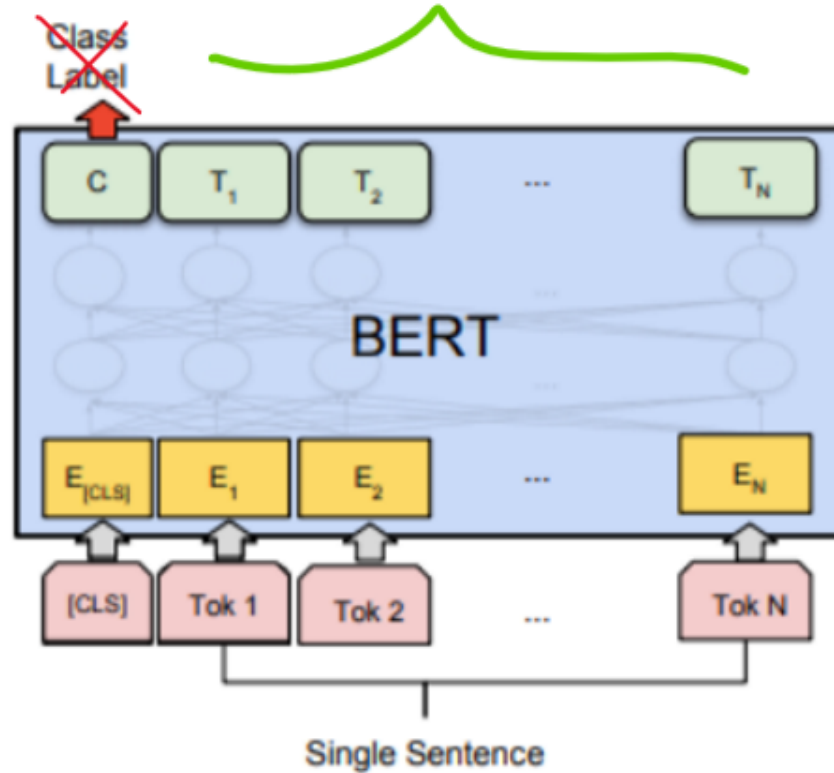
The fine-tuning approach isn't the only way to use BERT. Just like ELMo, you can use the pre-trained BERT to create contextualized word embeddings. Then you can feed these embeddings to your existing model – a process the paper shows yield results not far behind fine-tuning BERT on a task such as named-entity recognition.



Which vector works best as a contextualized embedding? I would think it depends on the task. The paper examines six choices (Compared to the fine-tuned model which achieved a score of 96.4):

```
from IPython.display import Image
Image(filename='bert.png')
```

Using These feature as word Embedding for each token and average it to find the embedding of each sentence



### 6.4.1. Finding AVG BERT Embedding

```
# Loading pretrained Model and Tokenizer
```

```
bert_model = BertModel.from_pretrained('bert-base-uncased')
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

bert_model.cuda()
print("Model loaded on gpu.." )
```

Model loaded on gpu..

```
# function to find embedding of test using Bert
```

```
def bert_embedding_of_text(text,max_len = 128):
```

```
    # Step 1: Tokenize
```

```
    tokens = tokenizer.tokenize(text)
```

```
    # Step 2: Add [CLS] and [SEP]
```

```
    tokens = ['[CLS]'] + tokens + ['[SEP]']
```

```
    # Step 3.1: Pad tokens
```

```
    if len(tokens) < max_len:
        padded_tokens = tokens + ['[PAD]'] * (max_len - len(tokens))
    else:
```

```

        padded_tokens = tokens[:max_len-1] + ['[SEP]']

    # Step 3.2: attention mask
    attn_mask = [1 if token != '[PAD]' else 0 for token in padded_tokens]

    # Step 4: Segment ids
    seg_ids = [0 for _ in range(len(padded_tokens))] #Optional!

    # Step 5: Get BERT vocabulary index for each token
    token_ids = tokenizer.convert_tokens_to_ids(padded_tokens)

    ## Bert Embeddings of words except [cls] token
    # Convert to pytorch tensors
    token_ids = torch.tensor(token_ids).unsqueeze(0).cuda()
    attn_mask = torch.tensor(attn_mask).unsqueeze(0).cuda()
    seg_ids = torch.tensor(seg_ids).unsqueeze(0).cuda()

    # Feed them to bert
    hidden_reps, cls_head = bert_model(token_ids, attention_mask = attn_mask,\
                                       token_type_ids = seg_ids)

    # Covert torch tensor to numpy array
    hidden_reps = hidden_reps.detach()

    return hidden_reps.squeeze()

#=====

# Function to find similarity using BERT avg Embedding using only 512 as max_length
def similarity_using_bert(text1, text2,max_length = 512):

    temp1 = bert_embedding_of_text(text1, max_length)
    temp2 = bert_embedding_of_text(text2, max_length)

    cos_sim = torch.nn.functional.cosine_similarity(temp1.mean(dim=1).unsqueeze(0),temp2.mean(dim=1).unsqueeze(0))

    return cos_sim.cpu().detach().numpy()[0]

```

## 6.4.2. Finding Similarity

```

# Finding similarity

similarity_score_bert = []
for i in tqdm(range(data.shape[0])):

    sentence1 = data['text1'].iloc[i]
    sentence2 = data['text2'].iloc[i]
    similarity_score_bert.append(similarity_using_bert(sentence1, sentence2, max_length = 512))

# Scaling the similarity between [0-1]
similarity_score_bert = minmax_scale(similarity_score_bert, feature_range=(0, 1

```

```
# Round off upto 2 decimal
similarity_score_bert = np.round(similarity_score_bert, decimals=3)
```

```
# Checking some similarity text
print(f"Index where similarity between text1 and text2 is greater than 0.9: {n
p.where(similarity_score_bert>0.8)[0]}")

print(f"Similarity Score of those index: {similarity_score_bert[np.where(simila
rity_score_bert>0.8)]}")
```

```
# Sample of text where similarity score is greater than 0.90
i = 3403

print("Sample of text where similarity score is greater than 0.9\n\n")
print(f"text1: \n{'-'*7}\n{data.iloc[i].text1}\n")
print(f"text2: \n{'-'*7}\n{data.iloc[i].text2}\n")
```

```
text1:
-----

holmes starts 2005 with gb events kelly holmes will start 2005 with a series of races in britain. holmes will make her first track appearance on home soil since winning double olympic gold in january s norwich union international in glasgow. she will also run in the grand prix in birmingham in february and may defend her indoor aaa 800m title in sheffield earlier this month. i am still competitive and still want to win she said. i m an athlete and i can t wait to get back on the track. she added: these events are also a great opportunity to thank the british public for the enormous levels of support they have given me from the moment i stepped off that plane from greece. the glasgow meeting will see holmes compete over 1500m in a five-way match against sweden france russia and italy.
```

```
text2:
-----

holmes starts 2005 with gb events kelly holmes will start 2005 with a series of races in britain. holmes will make her first track appearance on home soil since winning double olympic gold in january s norwich union international in glasgow. she will also run in the grand prix in birmingham in february and may defend her indoor aaa 800m title in sheffield earlier this month. i am still competitive and still want to win she said. i m an athlete and i can t wait to get back on the track. she added: these events are also a great opportunity to thank the british public for the enormous levels of support they have given me from the moment i stepped off that plane from greece. the glasgow meeting will see holmes compete over 1500m in a five-way match against sweden france russia and italy.
```

## 7.1. Saving Result

```

# Saving submission File
submission_w2v = pd.DataFrame({'Unique_ID':data.Unique_ID, 'similarity_score_w2v':similarity_score_w2v})
submission_w2v.to_csv(path_or_buf= 'submission_w2v.csv', sep=',',)

# Saving submission File
submission_tfidf_w2v = pd.DataFrame({'Unique_ID':data.Unique_ID, 'similarity_score_tfidf_w2v':similarity_score_tfidf_w2v})
submission_tfidf_w2v.to_csv(path_or_buf= 'submission_tfidf_w2v.csv', sep=',',)

# Saving submission File
submission_spacy = pd.DataFrame({'Unique_ID':data.Unique_ID, 'similarity_score_spacy':similarity_score_spacy})
submission_spacy.to_csv(path_or_buf= 'submission_spacy.csv', sep=',',)

# Saving submission File
submission_bert = pd.DataFrame({'Unique_ID':data.Unique_ID, 'similarity_score_bert':similarity_score_spacy})
submission_bert.to_csv(path_or_buf= 'submission_bert.csv', sep=',',)

```

## Conclusion

- These approaches are just some first cut solution to the problem.
- We can see using mentioned 4 approaches we are getting quite decent results; some approaches are obviously better than others because of better embeddings.
- TfIDF avg W2V has done really well to find the similarity scores.
- Results could be definitely improved using some addition techines like dimension reduction techniques and matrix factorisation techniques.
- Need to explore more "how to tackle unsupervised problem ? "

END :)